

CCES – College of Computer Engineering and Sciences

Prince Sattam Bin Abdulaziz University

CS600: Advanced AI

Introduction to Artificial Intelligence

Chapter 1 — Agents, Reasoning, and Learning

Dr. Khalil Chebil

Artificial Intelligence Course

What is Artificial Intelligence?
How to Achieve AI
Intelligent Agents

Learning Paradigms
Summary

Learning Objectives

- ▶ Distinguish **deductive reasoning** from **inductive learning**
- ▶ Understand the concept of **agents** and **environments**
- ▶ Explore **supervised**, **unsupervised**, and **reinforcement** learning
- ▶ Reflect on the **Turing Test** and the goal of AGI




Section 1

What is Artificial Intelligence?

What is Artificial Intelligence? | A working definition

Definition


A branch of computer science concerned with the **automation of intelligent behaviors**.




The Evolution of AI | Six defining phases

1. **Early Excitement** (1950s–60s)
optimism about reasoning machines
2. **First AI Winter** (1970s)
limited computational power
3. **Expert Systems Era** (1980s)
knowledge-based systems
4. **Second AI Winter** (late 80s–90s)
LISP market collapse, funding cuts
5. **ML Renaissance** (2000s–10s)
statistical methods
6. **Deep Learning Revolution** (2010s–)
neural nets + massive data

Roots of Artificial Intelligence | An interdisciplinary field




Timeline of AI History | Milestones



The Turing Test | Can a machine think? (Turing, 1950)

Procedure

1. All communication is text-based
2. A judge questions two hidden respondents
3. One is human, one is a computer
4. The machine *passes* if the judge cannot reliably tell them apart




Criticisms

Tests *deception* rather than intelligence; linguistic-only; needs no understanding or consciousness; anthropocentric.


Status

No system has convincingly passed the *full* test. AGI remains aspirational, not a present reality.

Subfields of AI | The AI landscape



- ▶ **Anomaly detection** in processes and equipment
- ▶ **Optimize processes** to improve yield
- ▶ **Smarter decisions** that minimize risk
- ▶ **Predict future scenarios** with neural networks



AI can be dangerous if **misused** or **poorly designed**. Key risks include:

- ▶ Job displacement
- ▶ Privacy concerns
- ▶ Bias and discrimination
- ▶ Autonomous weapons



!

Safety
Fairness
Accountability
Governance

Takeaway

Ethical development, transparency, and regulation are essential.

Section 2

How to Achieve AI

1. Phenomenal Approach

- ▶ Knowledge Representation
- ▶ Expert Systems & Planning
- ▶ Search
 - Uninformed: BFS, DFS
 - Informed: A*, Greedy

Deductive reasoning

2. Biological Approach

- ▶ Neural Networks (ANN)
- ▶ Evolutionary Algorithms
- ▶ Learning: supervised, unsupervised, reinforcement
- ▶ Agents: NLP, vision, robotics

Inductive learning

Deductive vs. Inductive Reasoning | The two schools of thought

Deductive Reasoning

Start from a **knowledge base** of facts; apply logic to make inferences.

Rules + Facts $\Rightarrow$ Conclusions

Search and logic-based methods.

Inductive Reasoning

Start from **observations**; generalize to broader rules or patterns.

Examples $\Rightarrow$ General Rule

The foundation of machine learning.

The greatest strength of deductive reasoning — interpretability — is also its greatest limitation.

► Constraint Satisfaction (CSP)

variables, domains, constraints

► Search Algorithms

BFS, DFS, UCS, A*, local search

► Logic-Based Systems

propositional & first-order logic, resolution

Expert Systems

► **Knowledge base:** facts + rules


► **Inference engine:** applies rules

► **User interface**

Forward chaining (data-driven) ■ Backward chaining (goal-driven)

Facts about John

- ▶ Running a temperature
- ▶ Coughing
- ▶ Colored phlegm



In practice, thousands of rules and cases are needed for a system to work well.

Section 3

Intelligent Agents

The Agent Abstraction | Perceive — Decide — Act


An **agent**:

1. **Perceives** the environment via *sensors*
2. **Acts** on it via *actuators*
3. Makes **decisions** to achieve goals

Formal definition

An agent function maps percept histories to actions:

$$f : P^* \rightarrow A$$




PEAS Description | Characterizing an agent

Agent	Performance	Environment	Actuators	Sensors
Cleaning robot	Cleanliness, battery	Room, dirt	Wheels, vacuum	Camera, dirt sensor
Chess program	Win/loss/draw	Board, opponent	Move selection	Board position
Self-driving car	Safety, speed	Roads, traffic	Steering, throttle	Cameras, LIDAR, GPS
Medical diagnosis	Patient health	Symptoms	Treatment plan	Lab tests, imaging
Chatbot	User satisfaction	Conversation	Text response	User input text

Performance ▪ Environment ▪ Actuators ▪ Sensors

Types of Agents | From reflexes to learning




Thermostat → simple reflex **Vacuum w/ memory** → model-based
GPS navigation → goal-based **Trading bot** → utility-based **AlphaGo** → learning

The Learning Agent | Four interacting components

- ▶ **Performance Element** — selects actions
- ▶ **Learning Element** — improves the performance element
- ▶ **Critic** — provides feedback on performance
- ▶ **Problem Generator** — suggests exploratory actions

Example

AlphaGo learned to master Go through self-play.



Types of Environments | Design depends on the world

Dimension	Type 1	Type 2	Example
Observability	Fully observable	Partially observable	<i>Chess vs. Poker</i>
Determinism	Deterministic	Stochastic	<i>Chess vs. Backgammon</i>
Episodes	Episodic	Sequential	<i>Image class. vs. Chess</i>
Dynamics	Static	Dynamic	<i>Crossword vs. Driving</i>
State space	Discrete	Continuous	<i>Chess vs. Robot control</i>
Agents	Single-agent	Multi-agent	<i>Puzzle vs. Poker</i>


Simplest: fully observable, deterministic, episodic, static, discrete, single-agent. Hardest: real-world robotics (the opposite of every dimension).



Section 4

Learning Paradigms

Three Types of Learning | Inductive learning from data



Supervised Learning | Learning with a teacher

Learn a function from **labeled** examples.

- ▶ Training data: $(X_1, y_1), \dots, (X_n, y_n)$
- ▶ Learn f such that $y_i \approx f(X_i)$
- ▶ Minimize loss: $L = \sum_{i=1}^n (y_i - f(X_i))^2$

Two tasks


- ▶ **Classification**: categorical (spam/not spam)
- ▶ **Regression**: numerical (house prices)

Example

Data $(1, 3), (2, 5), (3, 7), (4, 9)$

$$y = 2x + 1$$

Predict $x = 5 \Rightarrow y = 11$.




Unsupervised Learning | Discovering hidden structure

Find structure in **unlabeled** data.

- ▶ **Clustering**: K-means, hierarchical, DBSCAN
- ▶ **Dimensionality reduction**: PCA, t-SNE
- ▶ **Anomaly detection**: fraud, intrusion

K-Means (sketch)

1. Initialize k random centroids
2. Repeat until convergence:
 - a. Assign each point to nearest centroid
 - b. Update centroids to cluster means




Reinforcement Learning | Learning from interaction

An **agent** learns which actions to take by observing **rewards**.

- ▶ State s , action a , reward r , policy $\pi(s)$
- ▶ Goal: maximize cumulative reward

$$G_t = \sum_{k=0}^{\infty} \gamma^k r_{t+k}, \quad \gamma \in [0, 1]$$

Algorithms: Q-Learning, DQN (value-based); Policy Gradient, Actor-Critic (policy-based); AlphaZero (model-based).



Applications

AlphaGo, robotics, autonomous vehicles, data-center cooling.

Comparing Learning Types | At a glance

Aspect	Supervised	Unsupervised	Reinforcement
Data	Labeled examples	Unlabeled data	Interaction w/ environment
Feedback	Correct answers	None (implicit)	Rewards / penalties
Goal	Predict outputs	Discover patterns	Maximize reward
Training	(X, y) pairs	X only	(s, a, r, s') tuples
Example	Spam detection	Segmentation	Game playing



Section 5 Summary

Reasoning

- ▶ **Deductive:** rules + facts $\Rightarrow$ conclusions
- ▶ **Inductive:** examples $\Rightarrow$ general rules

Agents

Perceive–decide–act; PEAS; five agent types; environment dimensions.

Learning

- ▶ Supervised — labeled data
- ▶ Unsupervised — hidden structure
- ▶ Reinforcement — reward-driven

Looking ahead

Ch. 2: Search ■ Ch. 4–5: Logic ■ Ch. 6–10: ML.

Thank you!

Questions & Discussion

Next: Chapter 2 — Searching State Spaces